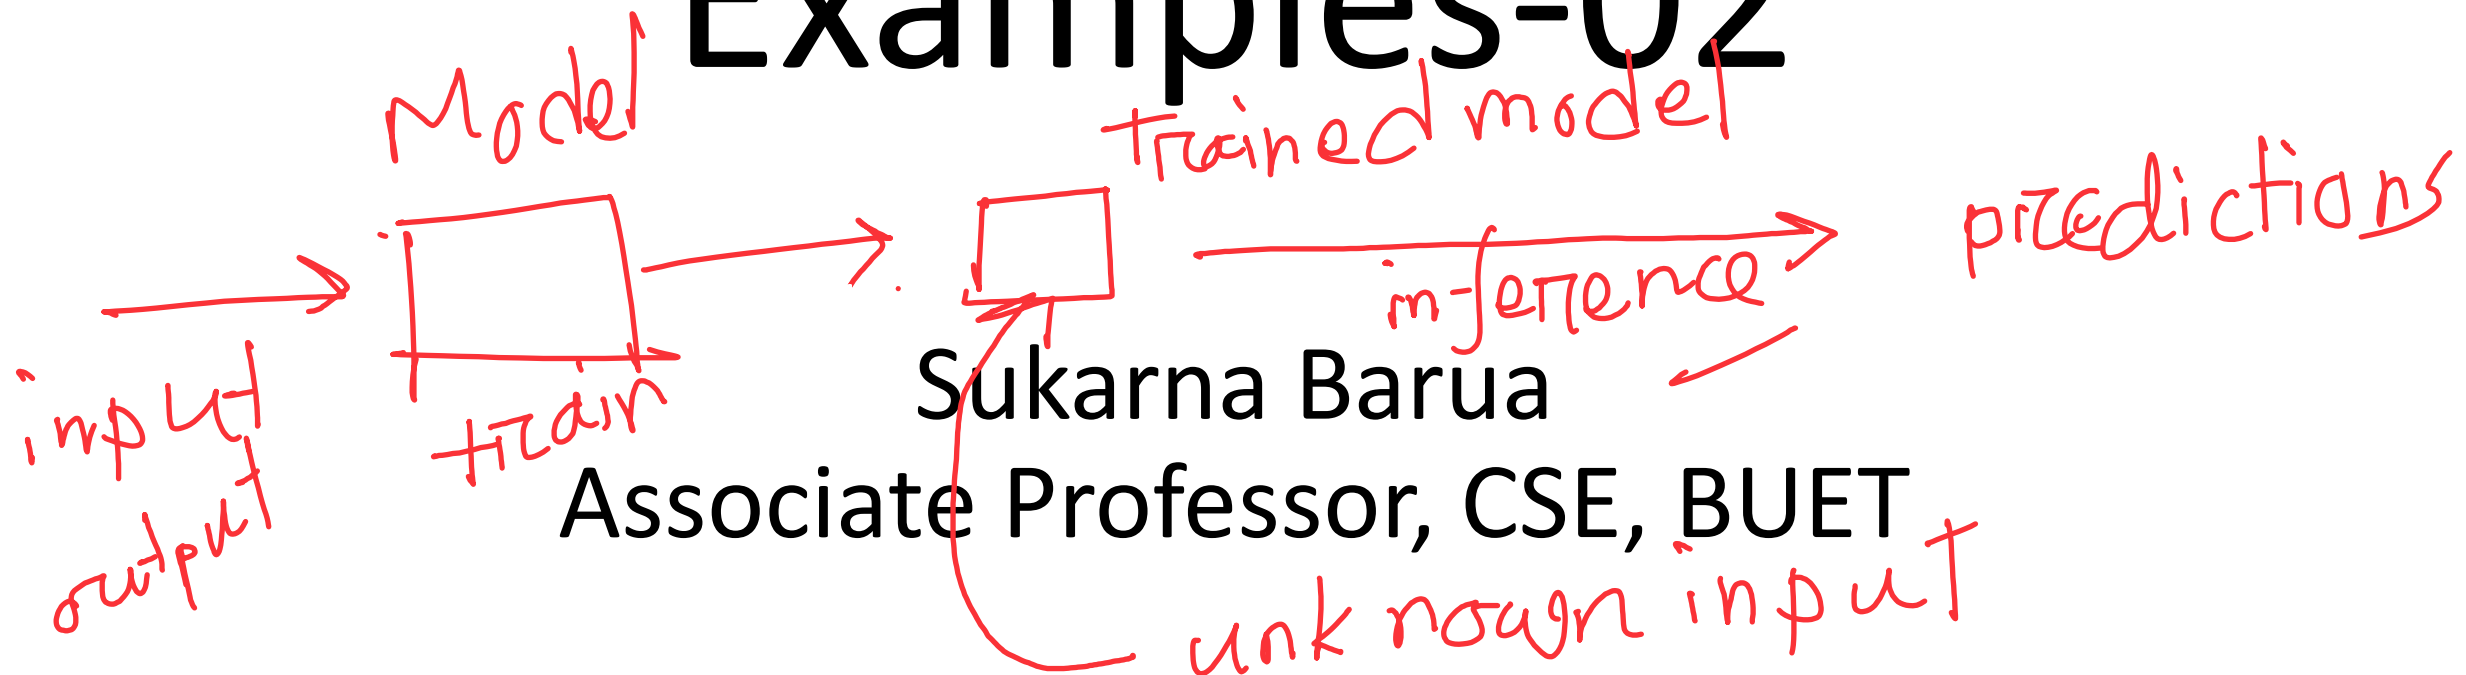


# Chapter 18 (AIAMA)

## Learning From Examples-02





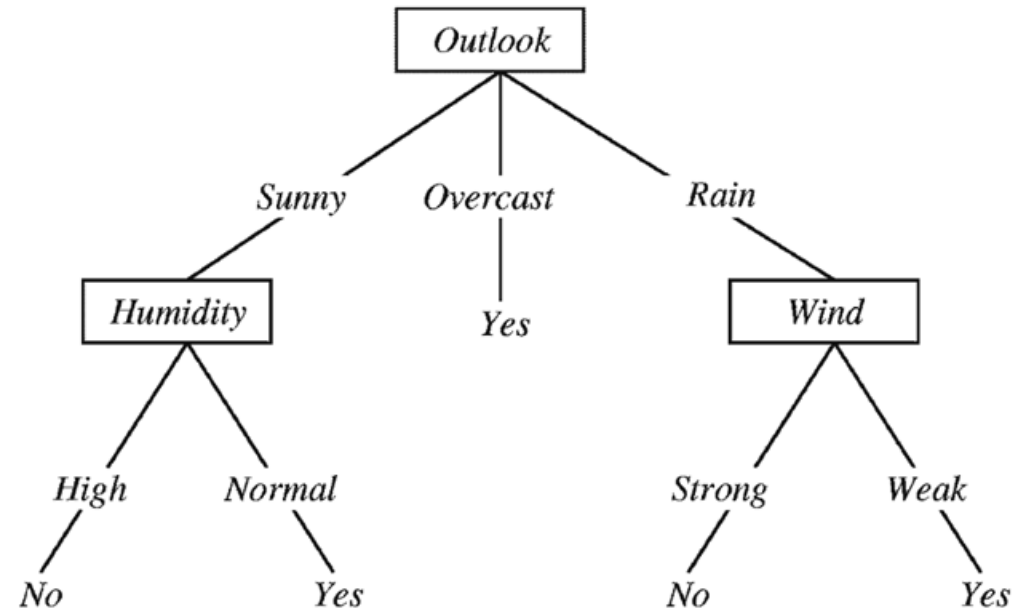
# Decision Tree

↓ Interpretable

- **A decision tree:**
  - Represents a function that takes as input a vector of attribute values
  - Returns a “decision”—a single output value.
  - The input and output values can be discrete or continuous.
  - For now we will concentrate on problems where
    - Inputs have discrete values
    - Output has exactly two possible values; this is Boolean classification, where each example input will be classified as *true* (a **positive** example) or *false* (a **negative** example).

# Decision Tree

- A decision tree for rain forecasting.
  - Input: {Outlook=Sunny, Humidity=Normal, Wind=Strong}
  - Output: Yes [Rainy]



# Decision Tree

- **In a decision tree:**

- Each **internal node** in the tree corresponds to a test of the value of one of the input attributes,  $A_i$
- The **branches** from the node are labeled with the possible values of the attribute,  $A_i = v_{ik}$ .
- Each **leaf node** in the tree is marked with a outcome to be returned by the function.

# Decision Tree: Example

- **Problem:** Build a decision tree to decide whether to wait at a restaurant.
- **Learning Goal:**
  - WillWait - to decide whether to wait for a table at a restaurant.
  - Goal is binary valued (i.e., binary classification task)
    - Values: {Yes, No}

# Decision Tree: Example

- **Input Attributes**

1. *Alternate*: whether there is a suitable alternative restaurant nearby.
2. *Bar* : whether the restaurant has a comfortable bar area to wait in.
3. *Fri/Sat*: *true* on Fridays and Saturdays, *false* otherwise.
4. *Hungry*: whether we are hungry.
5. *Patrons*: how many people are in the restaurant (values are None, Some, and Full ).
6. *Price*: the restaurant's price range (\$, \$\$, \$\$\$).

# Decision Tree: Example

- **Input Attributes**

- 7. *Raining*: whether it is raining outside.

- 8. *Reservation*: whether we made a reservation.

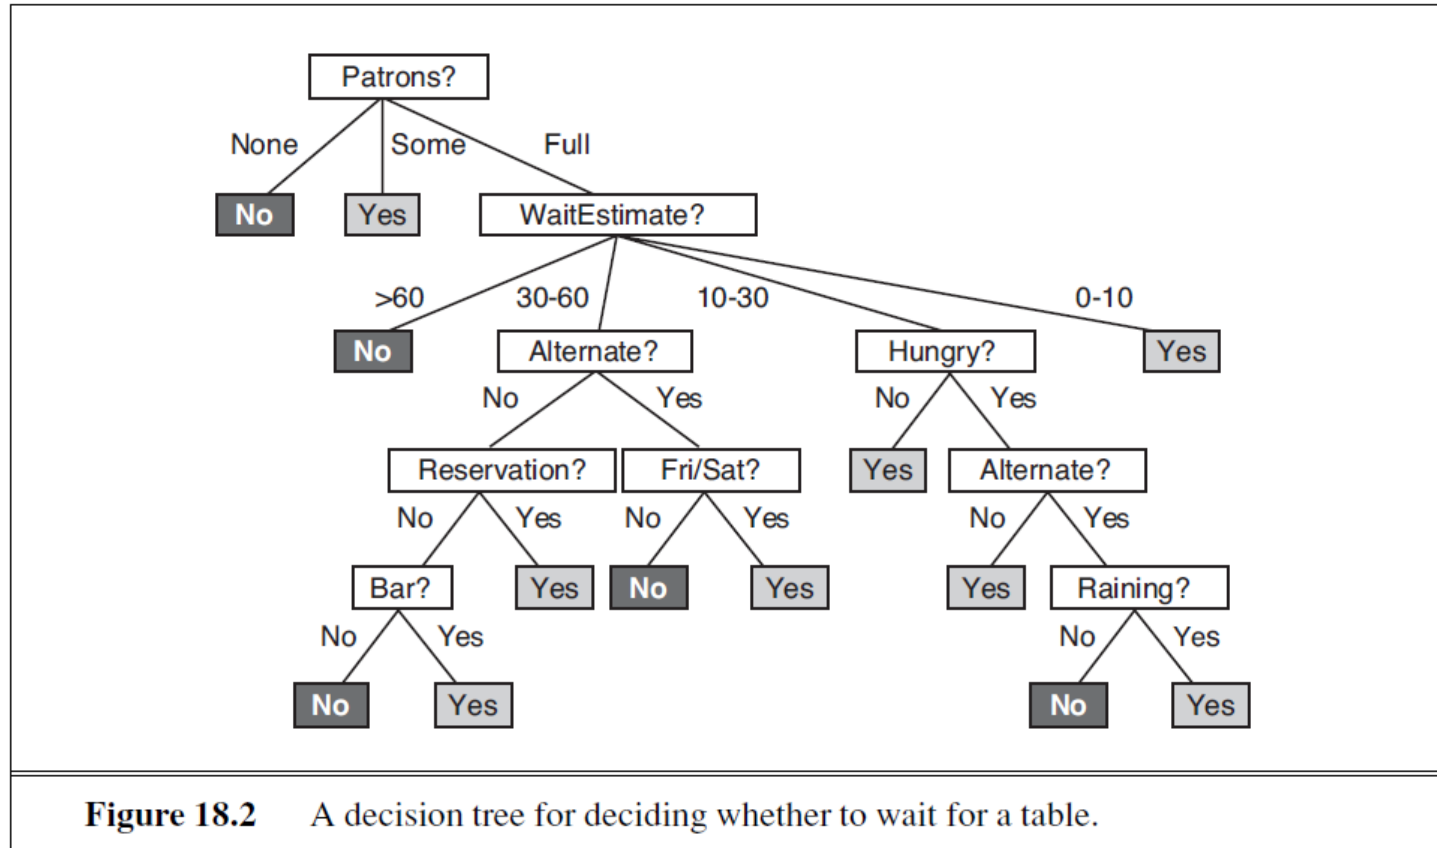
- 9. *Type*: the kind of restaurant (French, Italian, Thai, or burger).

- 10. *WaitEstimate*: the wait estimated by the host (0–10 minutes, 10–30, 30–60, or >60).

- Note that every variable has a small set of possible values; the value of *WaitEstimate*, for example, is not an integer, rather it is one of the four discrete values 0–10, 10–30, 30–60, or >60.

# Decision Tree: Example

- An example decision tree for the restaurant problem





# Expressiveness of Decision Trees

- In a Boolean decision tree, the goal attribute is true if and only if the input attributes satisfy one of the paths leading to a leaf with value true.
- $Goal \Leftrightarrow (Path1 \vee Path2 \vee \dots)$ , where each  $Path$  is a conjunction of attribute-value tests required to follow that path.
- Thus, the whole expression is equivalent to *disjunctive normal form*, which means that any function in propositional logic can be expressed as a decision tree.

# Expressiveness of Decision Trees

- Any Boolean function can be represented by a decision tree.
  - Consider the following Boolean function:  $f(A, B) = A + B$ .
  - Draw a decision tree for the function:

# Expressiveness of Decision Trees

- For a wide variety of problems, the decision tree format yields a nice, concise result.
- But some functions cannot be represented concisely:
  - For example, the majority function, which returns true if and only if more than half of the inputs are true, requires an exponentially large decision tree.
- Decision trees are good for some kinds of functions and bad for others.

# Expressiveness of Decision Trees

- **How many different decision trees can be obtained for a Boolean function with  $n$  variables?**
  - A truth table over  $n$  attributes has  $2^n$  rows, one for each combination of values of the attributes.
  - There are  $2^{2^n}$  different functions.
  - Each function can be represented by a decision tree. Hence, number of decision trees at least  $2^{2^n}$  [why more?]

# Expressiveness of Decision Trees

- **How many different decision trees can be obtained for a Boolean function with  $n$  variables?**
  - For 10 Boolean attributes of our restaurant problem there are  $2^{1024}$  or about  $10^{308}$  different functions to choose from.
  - Number of possible decision trees  $\geq 10^{308}$

# Inducing Decision Trees

- **Training data:** A set of examples where each example a pair  $(\underline{x}, y)$  pair where  $\underline{x}$  is a vector of values for the input attributes,  $y$  is a single Boolean output value.


Example	Input Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
$x_1$	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>0-10</i>	$y_1 = \text{Yes}$
$x_2$	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>30-60</i>	$y_2 = \text{No}$
$x_3$	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Some</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>0-10</i>	$y_3 = \text{Yes}$
$x_4$	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Thai</i>	<i>10-30</i>	$y_4 = \text{Yes}$
$x_5$	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>&gt;60</i>	$y_5 = \text{No}$
$x_6$	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Italian</i>	<i>0-10</i>	$y_6 = \text{Yes}$
$x_7$	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>0-10</i>	$y_7 = \text{No}$
$x_8$	<i>No</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Thai</i>	<i>0-10</i>	$y_8 = \text{Yes}$
$x_9$	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>&gt;60</i>	$y_9 = \text{No}$
$x_{10}$	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>Italian</i>	<i>10-30</i>	$y_{10} = \text{No}$
$x_{11}$	<i>No</i>	<i>No</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>0-10</i>	$y_{11} = \text{No}$
$x_{12}$	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>30-60</i>	$y_{12} = \text{Yes}$

**Figure 18.3** Examples for the restaurant domain.

# Inducing Decision Trees

- ✓ Build a decision tree that is -
  - ✓ Consistent with the examples [*Not always expected though, generalization may suffer*]
  - ✓ Is as small as possible. [*Occam's razor*]

# Inducing Decision Trees

- **Question:** Can we always find a consistent decision tree given a set of training examples?.
- **Answer:** Yes.
  -  However, if two training examples  $(\mathbf{x}, y)$  and  $(\mathbf{x}', y')$  have different outputs but same input attributes ( $\mathbf{x} = \mathbf{x}'$ ), then a consistent decision tree is not possible.



# Inducing Decision Trees

- Unfortunately, it is an intractable problem to find the smallest consistent tree; there is no way to efficiently search through *hypotheses space*.
  - An NP-hard problem!
- What can we do if cannot find the smallest decision tree?
- **Solution:** Use heuristics to find a closest one.
  - With some simple heuristics, we can find a good approximate solution: a small (*but not the smallest*) consistent tree.
  - This is a greedy approach. [Remember what is *greedy*]

# Decision Tree Learning

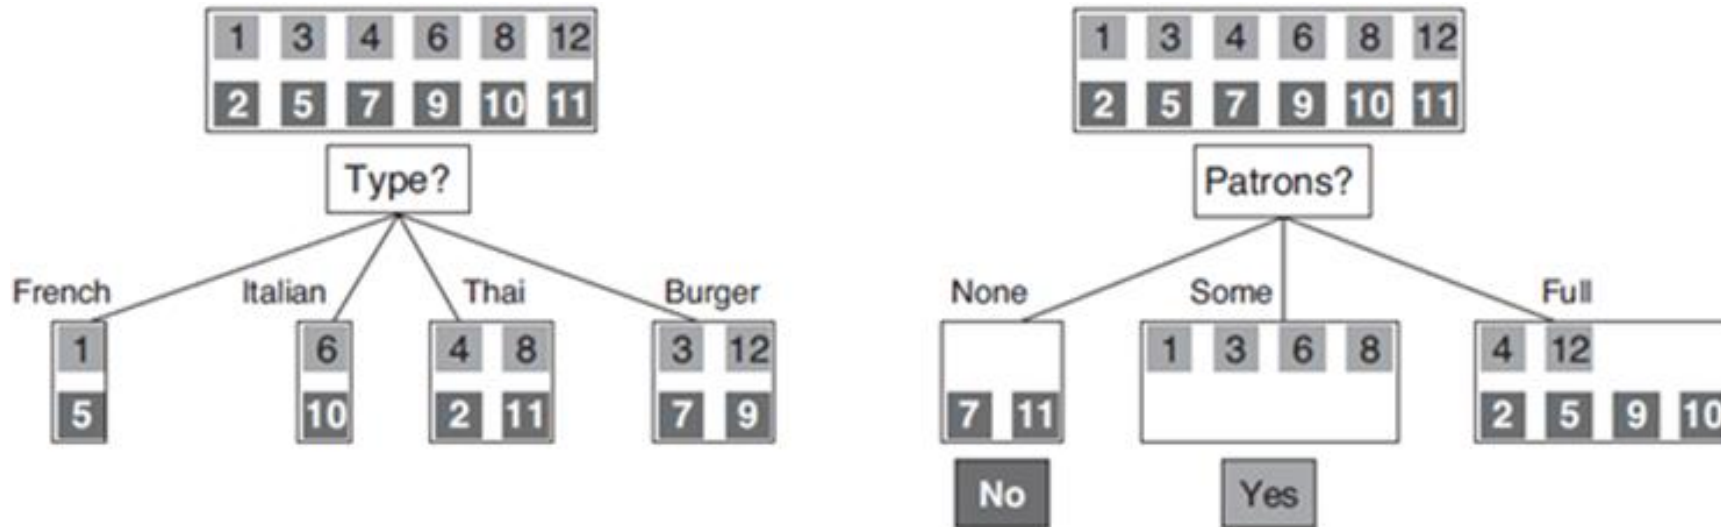
- **Greedy approach to build a decision tree:**
  - Start with empty decision tree.
  - Select an attribute to test at the next level [node]:
    - *Always select the most important attribute to test first. [greedy strategy]*
  - The test creates new branches and divides the problem into smaller subproblems.
  - Recurse on each child (created for each branch)

# Decision Tree Learning

- **Greedy strategy:** *Always select the most important attribute to test first.*
  - *Most important attribute* implies the one that makes the most difference to the classification of an example. [*Get leaves as early as possible*]
    - Get correct classification with a small number of test
    - All paths in the tree will be short
    - Tree as a whole will be shallow.
  - *Above greedy strategy is a local optimal choice, may not necessarily leads to the globally smallest tree!*

# Decision Tree Learning

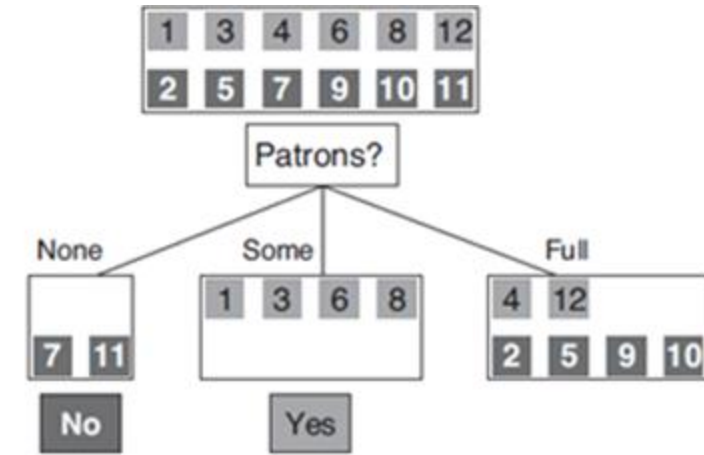
- Which attribute to test at root? Type vs Patron?
  - Type: all subsets (i.e., branches) needs further exploration.
  - Patron: two branches become leaves, only one needs further exploration.



# Decision Tree Learning Algorithm

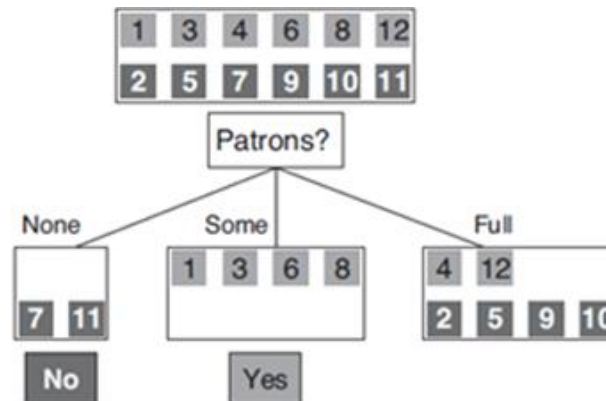
## ■ Decision tree construction:

- Step 1: Test an attribute at each node.
- Step 2: Partition the examples according to values and create child nodes with relevant examples.
- Step 3: Now consider child node for further tests of attributes except the one which have already been tested in the hierarchy (*recursive operation*)



# Decision Tree Learning Algorithm

- **Four cases to consider at a child node:**
  - Case 1: All examples are positive (or all negative) → we are done; we can answer Yes or No. [e.g., *Patron=None*]
  - Case 2: Some positive and some negative examples → then choose the best attribute to split them recursively. [e.g., *Patron=Full*]



# Decision Tree Learning Algorithm

- **Four cases to consider at a child node:**
  - Case 3: No examples left → No example has been observed for this combination of attribute values, and we return a default value (e.g., *plurality of parent node*)
  - Case 4: No attributes left, but both positive and negative examples remain → These examples have exactly the same description, but different classifications. This can happen because there is an error or **noise** in the data.
    - *In this case return the plurality classification of the remaining examples.*  
[tree will not be consistent]

# Decision Tree Learning Algorithm

- **Algorithm pseudocode**

```
function DECISION-TREE-LEARNING(examples, attributes, parent_examples) returns  
a tree  
  
  if examples is empty then return PLURALITY-VALUE(parent_examples)  
  else if all examples have the same classification then return the classification  
  else if attributes is empty then return PLURALITY-VALUE(examples)  
  else  
     $A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$   
    tree  $\leftarrow$  a new decision tree with root test A  
    for each value  $v_k$  of A do  
       $\text{exs} \leftarrow \{e : e \in \text{examples} \text{ and } e.A = v_k\}$   
      subtree  $\leftarrow$  DECISION-TREE-LEARNING(exs, attributes – A, examples)  
      add a branch to tree with label (A =  $v_k$ ) and subtree subtree  
  return tree
```

- Function PLURALITY-VALUE selects the most common class/output among the examples



# Choosing the Most Important Attribute

- **Which attribute is the most important now?**
  - **Perfect attribute:** One that splits into subsets where each subset contain either all positive or all negative examples. [*all branches become leaf nodes*]
  - **Useless attribute:** One that splits into subsets where each subset contain fairly equal mix of positive and negative examples. [*all branches need recursive exploration*]

# Choosing the Most Important Attribute

- **Perfect vs. useless? How to measure?**
  - A formal measure of perfect vs useless: **Entropy**
  - The fundamental quantity in information theory (*Shannon and Weaver, 1949*).
  - Entropy is a measure of the uncertainty of a random variable.
  - Acquisition of information corresponds to a reduction in entropy.
  - A random variable with only one value—a coin that always comes up heads—has no uncertainty and thus its entropy is defined as zero; thus, we gain no information by observing its value.

# Choosing the Most Important Attribute

- **Entropy:** Average number of bits per symbol to encode information.
- The roll of a fair *four*-sided die has 2 bits of entropy, because it takes two bits to describe one of four equally probable choices.
- An unfair coin that comes up heads 99% of the time.
  - This coin has less uncertainty than the fair coin—if we guess heads we'll be wrong only 1% of the time—it's entropy measure should be close to zero, but positive.

# Entropy Measure

- In general, the entropy  $H(V)$  of a random variable  $V$  with values  $v_k$ , each with probability  $P(v_k)$ , is defined as :

$$H(V) = \sum_k P(v_k) \log_2 \frac{1}{P(v_k)} = - \sum_k P(v_k) \log_2 P(v_k)$$

# Entropy Measure

- Verify that entropies measures are correct.

We can check that the entropy of a fair coin flip is indeed 1 bit:

$$H(Fair) = -(0.5 \log_2 0.5 + 0.5 \log_2 0.5) = 1 .$$

If the coin is loaded to give 99% heads, we get

$$H(Loaded) = -(0.99 \log_2 0.99 + 0.01 \log_2 0.01) \approx 0.08 \text{ bits.}$$

# Entropy Measure

- Define  $B(q)$  as the entropy of a Boolean random variable that is true with probability  $q$ :

$$B(q) = -(q \log_2 q + (1 - q) \log_2 (1 - q))$$

- If a training set contains  $p$  positive examples and  $n$  negative examples, then the entropy of the goal attribute on the whole set is:

$$H(Goal) = B\left(\frac{p}{p + n}\right) .$$

# Entropy Measure in Decision Tree

- **Decision tree contest:** Entropy represents an impurity measure of the set.
  - A set with 5 positive and 5 negative examples: Most impure, entropy should be highest.
  - A set with 10 positive and 0 negative examples: Purest, entropy should be the lowest.

# Entropy Before Split

- A set with  $p$  positive and  $n$  negative examples.

- Entropy of the set is:  $B\left(\frac{p}{p+n}\right)$



# Entropy After Splitting

- An attribute  $A$  with  $d$  distinct values divides the training set  $E$  into subsets  $E_1, E_2, \dots, E_d$ .
- Each subset  $E_k$  has  $p_k$  positive examples and  $n_k$  negative examples, with entropy of  $B(p_k/(p_k + n_k))$  bits of information. [ $E_k$  contains examples with  $A = kth$  value]
- A randomly chosen example from the training set has the  $kth$  value for the attribute with probability  $(p_k + n_k)/(p + n)$ , so the expected entropy (*weighted average*) of the  $d$  subsets after splitting on attribute  $A$  is:

$$Remainder(A) = \sum_{k=1}^d \frac{p_k + n_k}{p + n} B\left(\frac{p_k}{p_k + n_k}\right)$$

# Information Gain

- **Information Gain:** The amount of information gain is the amount of entropy reduction after the split on attribute A.
- Assume  $E_{BS}$  = Entropy before split,  $E_{AS}$  = Entropy after split [*weighted average*]
  - Hence, information gain ( = *reduction of entropy*):

$$Gain = E_{BS} - E_{AS} \text{ [reduction of entropy]}$$

- Hence, this is simply:

$$Gain(A) = B\left(\frac{p}{p+n}\right) - Remainder(A)$$

# Information Gain

- **Information Gain:**  $Gain(A) = B(\frac{p}{p+n}) - Remainder(A)$
- Compute the information gain for the attributes Patron and Type:

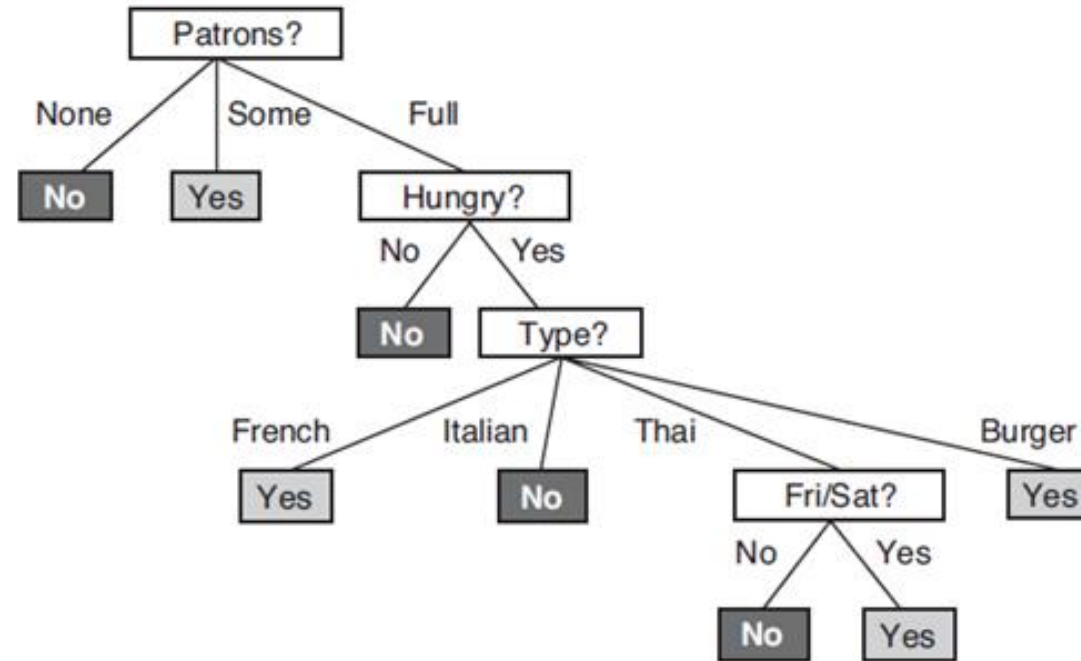
$$Gain(Patrons) = 1 - [\frac{2}{12}B(\frac{0}{2}) + \frac{4}{12}B(\frac{4}{4}) + \frac{6}{12}B(\frac{2}{6})] \approx 0.541 \text{ bits},$$

$$Gain(Type) = 1 - [\frac{2}{12}B(\frac{1}{2}) + \frac{2}{12}B(\frac{1}{2}) + \frac{4}{12}B(\frac{2}{4}) + \frac{4}{12}B(\frac{2}{4})] = 0 \text{ bits},$$

- Patron is a better attribute than Gain! Hence, choose Patron over Type.
- *Choose the attribute which gives the highest information gain!*

# Final Decision Tree

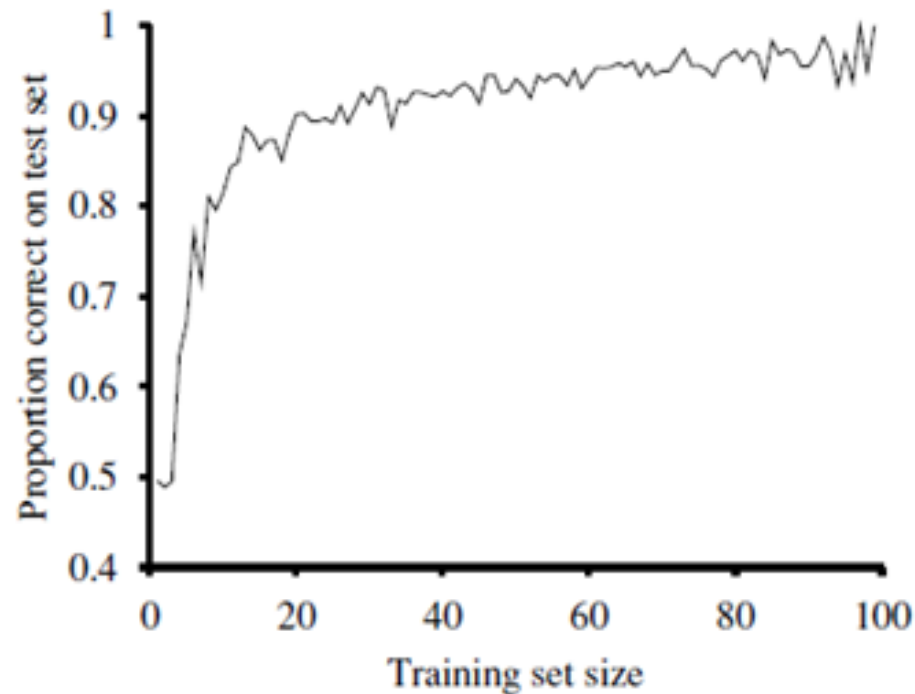
- Final decision tree constructed from given examples.



- Question:** What is the information gain for Type attribute at level 3?

# Decision Tree Learning Curve

- Experiment with 100 examples: Construct decision trees
  - Split into train and test (e.g., 1 and 99, 2 and 98, etc.)
  - Random split 20 times and report average accuracy on test set
  - Note: As the training size grows, accuracy increases.



# Overfitting

- **Overfitting:** Too much importance on every training example
  - Complex tree
  - Suffers generalization: very low training error, but very high test error.
  - Misses important concepts
- A consistent decision tree over training data may result in a complex tree
  - Noisy data may also induce complexity in the tree

# Overfitting Solution

- **Pre-prune:** Prune the tree before it gets large
  - Early stopping: limit depth during tree construction
- **Post-pruning:** Prune the tree after construction
  - Remove irrelevant nodes
  - Replace internal nodes with most common class [*Only replace if test error do not increase*]

# Decision Tree Issues

- **Missing values:** Some examples have missing value in some attributes
  - Solution: Replace missing value with mean of the attribute over the entire dataset
- **Continuous-valued attribute:** For example, height is continuous-valued.
  - Convert to categorical attribute:
    - Height > 40cm: *Tall* and Height <40cm: *Short*



8. (a) Consider the following data set comprised of three binary input attributes ( $A_1, A_2$ , and  $A_3$ ) and one binary output  $y$ :

(1)

Example	$A_1$	$A_2$	$A_3$	Output $y$
$x_1$	1	0	0	0
$x_2$	1	0	1	0
$x_3$	0	1	0	0
$x_4$	1	1	1	1
$x_5$	1	1	0	1

Construct a decision tree for these data. Show the entropy and information-gain computations made to determine the attribute to split at each node.

Suppose a bank wants to decide whether to approve a small loan (output  $y$ ). They use three simple binary attributes about an applicant:

- $A_1$  = Has a stable job? (1 = Yes, 0 = No)
- $A_2$  = Has no unpaid loans? (1 = Yes, 0 = No)
- $A_3$  = Owns a house? (1 = Yes, 0 = No)
- Output  $y$  = Loan approved? (1 = Yes, 0 = No)

train the  
dataset

$A_1 A_2 A_3$